

# 1 Business Understanding

## 1.1 Project Overview

This project outlines analysis based on movie datasets for business-stakeholders new to the movie industry.

## 1.2 Business problem:

Your company now sees all the big companies creating original video content and they want to get in on the fun. They have decided to create a new movie studio, but they don't know anything about creating movies. You are charged with exploring what types of films are currently doing the best at the box office. You must then translate those findings into actionable insights that the head of your company's new movie studio can use to help decide what type of films to create.

## 1.3 Project objectives:

### 1.3.1 Main Objective

To analyze movie data and uncover patterns in sales, popularity, ratings, and director influence across genres, providing actionable insights for business growth and strategy.

### 1.3.2 Specific Objectives

1. Identify which genres generate the most revenue and analyze trends contributing to their sales performance.
2. Understand which genres are most popular among audiences and explore factors driving their popularity.
3. Examine the ratings of top genres to evaluate their critical reception.
4. Identify top selling movies with their related genre.
5. Determine most popular directors
6. Identify the top selling directors
7. Identify top popular movie ratings

### 1.3.3 The Data

We used the folder `zippedData` that are movie datasets from the following websites:

- [Box Office Mojo \(https://www.boxofficemojo.com/\)](https://www.boxofficemojo.com/)
- [IMDB \(https://www.imdb.com/\)](https://www.imdb.com/)
- [Rotten Tomatoes \(https://www.rottentomatoes.com/\)](https://www.rottentomatoes.com/)
- [TheMovieDB \(https://www.themoviedb.org/\)](https://www.themoviedb.org/)
- [The Numbers \(https://www.the-numbers.com/\)](https://www.the-numbers.com/)

Here are the datasets our crucial for analysis: `bom.movie_gross.csv` , `rt.movie_info.tsv` , `tmdb.movies.csv` a sql database: `im.db` where tables considered were; `movie_basics` , `movie_rating`

## 2 Data Understanding

Here will need to understand our data. This involves getting the relevant information from each dataset crucial for our analysis.

We start by loading the various datasets reviewing their various information based on the columns and check which information is necessary for our analysis before beginning the data cleaning.

```
In [1]: #importing libraries for data manipulation (pandas, numpy) and visualization
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
import sqlite3
import warnings
import numpy as np
import pandas as pd
from scipy.stats import norm
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Suppress all warnings
warnings.filterwarnings('ignore')
```

```
In [2]: # set the maximum number of columns to 40 to display all columns
pd.set_option('display.max_columns', 40)
```

Loading information from **rt.movie\_info.tsv** dataset

```
In [3]: #Loading the dataset and checking the top five columns
movie_df = pd.read_csv('Datasets/rt.movie_info.tsv', sep='\t')
movie_df.head()
```

Out[3]:

	id	synopsis	rating	genre	director	writer	theater_c
0	1	This gritty, fast-paced, and innovative police...	R	Action and Adventure Classics Drama	William Friedkin	Ernest Tidyman	Oct 9
1	3	New York City, not-too-distant-future: Eric Pa...	R	Drama Science Fiction and Fantasy	David Cronenberg	David Cronenberg Don DeLillo	Aug 17
2	5	Illeana Douglas delivers a superb performance ...	R	Drama Musical and Performing Arts	Allison Anders	Allison Anders	Sep 15
3	6	Michael Douglas runs afoul of a treacherous su...	R	Drama Mystery and Suspense	Barry Levinson	Paul Attanasio Michael Crichton	Dec 9
4	7	NaN	NR	Drama Romance	Rodney Bennett	Giles Cooper	



```
In [4]: #checking the shape getting information of the rows and columns
movie_df.shape
```

Out[4]: (1560, 12)

In [5]: *#checking information for each column of the dataset*  
 movie\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1560 entries, 0 to 1559
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                     1560 non-null   int64
1   synopsis               1498 non-null   object
2   rating                1557 non-null   object
3   genre                 1552 non-null   object
4   director              1361 non-null   object
5   writer                1111 non-null   object
6   theater_date          1201 non-null   object
7   dvd_date              1201 non-null   object
8   currency              340 non-null    object
9   box_office            340 non-null    object
10  runtime               1530 non-null   object
11  studio                494 non-null    object
dtypes: int64(1), object(11)
memory usage: 146.4+ KB
```

Loading the **bom.movie\_gross.csv**

In [6]: *#Loading the dataset*  
 gross\_df = pd.read\_csv("Datasets/bom.movie\_gross.csv")

In [7]: *#checking the top 5 columns*  
 gross\_df.head()

Out[7]:

	title	studio	domestic_gross	foreign_gross	year
0	Toy Story 3	BV	415000000.0	652000000	2010
1	Alice in Wonderland (2010)	BV	334200000.0	691300000	2010
2	Harry Potter and the Deathly Hallows Part 1	WB	296000000.0	664300000	2010
3	Inception	WB	292600000.0	535700000	2010
4	Shrek Forever After	P/DW	238700000.0	513900000	2010

In [8]: *#Viewing information for each column*  
gross\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3387 entries, 0 to 3386
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   title           3387 non-null   object
1   studio          3382 non-null   object
2   domestic_gross  3359 non-null   float64
3   foreign_gross   2037 non-null   object
4   year            3387 non-null   int64
dtypes: float64(1), int64(1), object(3)
memory usage: 132.4+ KB
```

In [9]: *#checking the statistical information for numerical columns*  
gross\_df.describe()

Out[9]:

	domestic_gross	year
count	3.359000e+03	3387.000000
mean	2.874585e+07	2013.958075
std	6.698250e+07	2.478141
min	1.000000e+02	2010.000000
25%	1.200000e+05	2012.000000
50%	1.400000e+06	2014.000000
75%	2.790000e+07	2016.000000
max	9.367000e+08	2018.000000

Loading the **tmdb.movies.csv** dataset

In [10]: *#Loading the dataset and viewing the first five columns*

```
tmdb_df = pd.read_csv("Datasets/tmdb.movies.csv")
tmdb_df.head()
```

Out[10]:

	Unnamed: 0	genre_ids	id	original_language	original_title	popularity	release_date
0	0	[12, 14, 10751]	12444	en	Harry Potter and the Deathly Hallows: Part 1	33.533	2011-11-18
1	1	[14, 12, 16, 10751]	10191	en	How to Train Your Dragon	28.734	2010-06-10
2	2	[12, 28, 878]	10138	en	Iron Man 2	28.515	2010-06-18
3	3	[16, 35, 10751]	862	en	Toy Story	28.005	1995-10-30
4	4	[28, 878, 12]	27205	en	Inception	27.920	2010-07-16

In [11]: *#Checking for information of each column from the dataset*

```
tmdb_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26517 entries, 0 to 26516
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            26517 non-null  int64
1   genre_ids             26517 non-null  object
2   id                    26517 non-null  int64
3   original_language     26517 non-null  object
4   original_title        26517 non-null  object
5   popularity            26517 non-null  float64
6   release_date          26517 non-null  object
7   title                 26517 non-null  object
8   vote_average          26517 non-null  float64
9   vote_count            26517 non-null  int64
dtypes: float64(2), int64(3), object(5)
memory usage: 2.0+ MB
```

```
In [12]: #Viewing the statistical information
tmdb_df.describe()
```

Out[12]:

	Unnamed: 0	id	popularity	vote_average	vote_count
count	26517.00000	26517.000000	26517.000000	26517.000000	26517.000000
mean	13258.00000	295050.153260	3.130912	5.991281	194.224837
std	7654.94288	153661.615648	4.355229	1.852946	960.961095
min	0.00000	27.000000	0.600000	0.000000	1.000000
25%	6629.00000	157851.000000	0.600000	5.000000	2.000000
50%	13258.00000	309581.000000	1.374000	6.000000	5.000000
75%	19887.00000	419542.000000	3.694000	7.000000	28.000000
max	26516.00000	608444.000000	80.773000	10.000000	22186.000000

Connecting to the SQL database

```
In [13]: #connecting to the sql database
conn = sqlite3.Connection('Datasets/im.db')
```

```
In [14]: #getting table names
cursor = conn.cursor()
cursor.execute("""SELECT name
                FROM sqlite_master
                WHERE type = 'table';""")
print(cursor.fetchall())


[('movie_basics',), ('directors',), ('known_for',), ('movie_akas',), ('movie_ratings',), ('persons',), ('principals',), ('writers',)]
```

```
In [15]: #Loading the movie basics table
mbasics_df = pd.read_sql("""SELECT * FROM movie_basics;""", conn)
```

In [16]: *#Checking the top five columns*  
 mbasics\_df.head()

Out[16]:

	movie_id	primary_title	original_title	start_year	runtime_minutes	
0	tt0063540	Sunghursh	Sunghursh	2013	175.0	Action,Crit
1	tt0066787	One Day Before the Rainy Season	Ashad Ka Ek Din	2019	114.0	Biograp
2	tt0069049	The Other Side of the Wind	The Other Side of the Wind	2018	122.0	
3	tt0069204	Sabse Bada Sukh	Sabse Bada Sukh	2018	NaN	Come
4	tt0100275	The Wandering Soap Opera	La Telenovela Errante	2017	80.0	Comedy,Dram



In [17]: *#Checking for the information of each column*  
 mbasics\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 146144 entries, 0 to 146143
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   movie_id              146144 non-null object
1   primary_title         146144 non-null object
2   original_title        146123 non-null object
3   start_year            146144 non-null int64
4   runtime_minutes       114405 non-null float64
5   genres                140736 non-null object
dtypes: float64(1), int64(1), object(4)
memory usage: 6.7+ MB
```

In [18]: *#Loading the the movie\_ratings table from the database*  
 rating\_df = pd.read\_sql("""SELECT \* FROM movie\_ratings;""",conn)

In [19]: *#Checking the top 5 columns*  
 rating\_df.head()

Out[19]:

	movie_id	averagerating	numvotes
0	tt10356526	8.3	31
1	tt10384606	8.9	559
2	tt1042974	6.4	20
3	tt1043726	4.2	50352
4	tt1060240	6.5	21



In [20]: *#Checking for the information from each column*  
 rating\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 73856 entries, 0 to 73855
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  -
0   movie_id        73856 non-null  object
1   averagerating   73856 non-null  float64
2   numvotes        73856 non-null  int64
dtypes: float64(1), int64(1), object(1)
memory usage: 1.7+ MB
```

In [21]: *#checking for the statistical information*  
 rating\_df.describe()

Out[21]:

	averagerating	numvotes
count	73856.000000	7.385600e+04
mean	6.332729	3.523662e+03
std	1.474978	3.029402e+04
min	1.000000	5.000000e+00
25%	5.500000	1.400000e+01
50%	6.500000	4.900000e+01
75%	7.400000	2.820000e+02
max	10.000000	1.841066e+06

In [22]: directors\_df = pd.read\_sql("""SELECT \* FROM directors;""",conn)

In [23]: directors\_df.head()

Out[23]:

	movie_id	person_id
0	tt0285252	nm0899854
1	tt0462036	nm1940585
2	tt0835418	nm0151540
3	tt0835418	nm0151540
4	tt0878654	nm0089502

```
In [24]: directors_df.tail()
```

```
Out[24]:
```

	movie_id	person_id
291169	tt8999974	nm10122357
291170	tt9001390	nm6711477
291171	tt9001494	nm10123242
291172	tt9001494	nm10123248
291173	tt9004986	nm4993825

```
In [25]: directors_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 291174 entries, 0 to 291173
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   movie_id    291174 non-null  object
 1   person_id   291174 non-null  object
dtypes: object(2)
memory usage: 4.4+ MB
```

```
In [26]: conn.close()
```

## 3 Data Preparation

### 3.1 Data Cleaning

After looking into our dataset and reviewing what categorical and numerical data we will require we begin our data cleaning process by:

1. Dropping columns unnecessary for our analysis
2. Checking for missing values
3. Rectifying column arrangement and uniformity
4. Checking for outliers
5. Dropping duplicates
6. Changing the data types

## 3.2 Cleaning for the movie\_df dataset

### 3.2.1 Dropping unnecessary columns

```
In [27]: # drop those columns with more than 1000 non-null rows
movie_df = movie_df.drop(['currency', 'studio', 'writer', 'synopsis', 'runtime'])
```

### 3.2.2 Checking for missing values

```
In [28]: #Checking for missing values in each column
movie_df.isna().sum()
```

```
Out[28]: id                0
rating                  3
genre                   8
director              199
theater_date          359
box_office            1220
dtype: int64
```

```
In [29]: #replacing movie genre nulls with mode and confirming the changes
genre_mode = movie_df.genre.mode()[0]
movie_df.genre.fillna(genre_mode, inplace=True)
movie_df.genre.isna().sum()
```

```
Out[29]: 0
```

```
In [30]: #replacing movie rating nulls with mode
rating_mode = movie_df.rating.mode()[0]
movie_df.rating.fillna(rating_mode, inplace=True)
movie_df.rating.isna().sum()
```

```
Out[30]: 0
```

```
In [31]: #drop the rest with nulls
movie_df.dropna(inplace=True)
```

```
In [32]: #Confirming that there are no null values
movie_df.isnull().sum().any()
```

```
Out[32]: False
```

```
In [33]: #Renaming columns
movie_df.rename(columns={'theater_date': 'year'}, inplace=True)
```

```
In [34]: #changing the box office values to numerical(box office sales a movie makes
#noted it was not numerical since it gave out an error when trying to fill na
# Step 1: Remove commas
movie_df['box_office'] = movie_df['box_office'].str.replace(',', '')

# Step 2: Convert to integers
movie_df['box_office'] = movie_df['box_office'].astype(int)
```

```
In [35]: #Confirming changes to the dataset
movie_df.head()
```

Out[35]:

	id	rating	genre	director	year	box_office
1	3	R	Drama Science Fiction and Fantasy	David Cronenberg	Aug 17, 2012	600000
6	10	PG-13	Comedy	Jake Kasdan	Jan 11, 2002	41032915
7	13	R	Drama	Ray Lawrence	Apr 27, 2006	224114
8	14	R	Drama	Taylor Hackford	Jun 30, 2010	134904
15	22	R	Comedy Drama Mystery and Suspense	George Hickenlooper	Dec 17, 2010	1039869

### 3.3 Cleaning for the gross\_df dataset

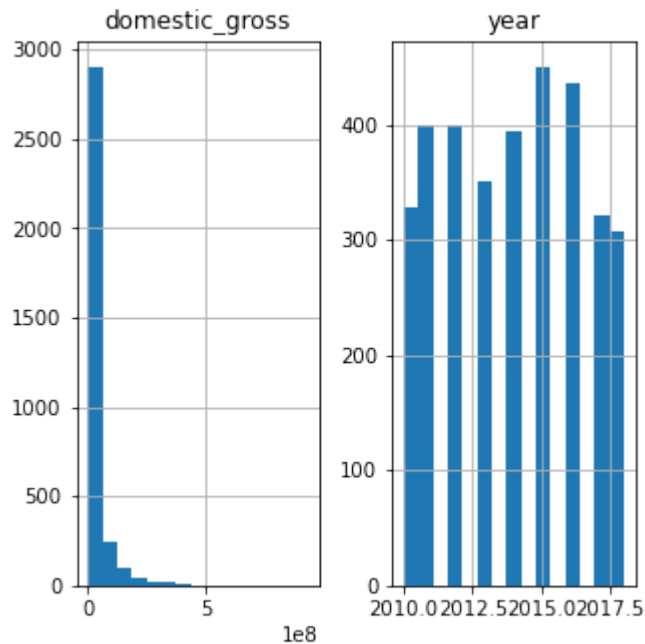
```
In [36]: #Dropping unnecessary columns
gross_df.drop(["studio"],axis=1,inplace=True)
```

#### 3.3.1 Checking for missing values

```
In [37]: #checking for any missng values
gross_df.isna().sum()
```

```
Out[37]: title                0
domestic_gross              28
foreign_gross             1350
year                        0
dtype: int64
```

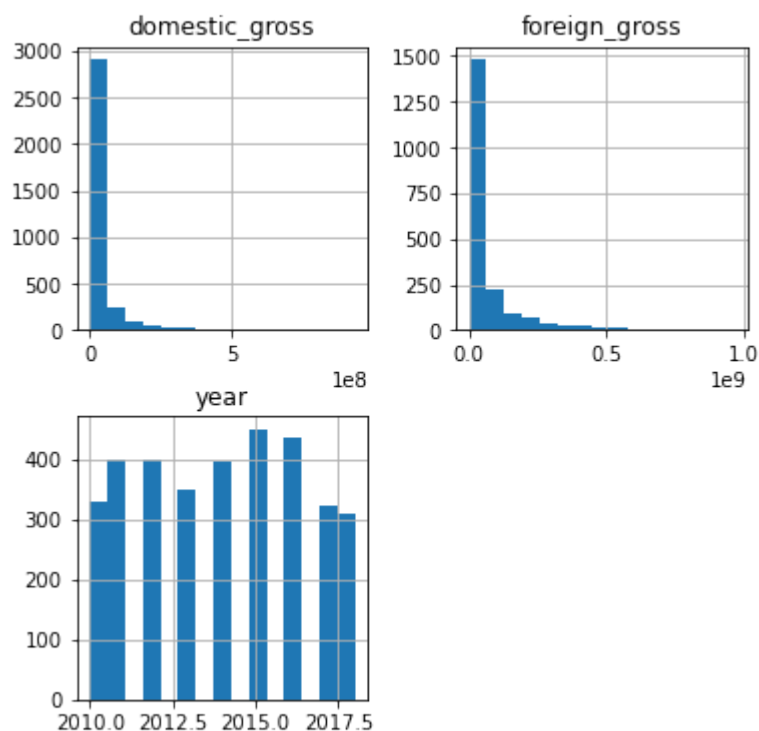
```
In [38]: #checking for the distribution for the numerical values
gross_df.hist(bins=15,figsize=(5,5));
```



From the histograms we noticed the foreign\_gross column is not numerical so we change its data type to float as shown below:

```
In [39]: # converting 'foreign_gross' to float
gross_df['foreign_gross'] = pd.to_numeric(gross_df['foreign_gross'],errors=
```

```
In [40]: #confirming the changes
gross_df.hist(bins=15,figsize=(6,6));
```



We notice that the distribution has skewness hence we use median to fill in the missing values

```
In [41]: #replacing gross for domestic and foreign with median
gross_df["domestic_gross"] = gross_df["domestic_gross"].fillna(gross_df["domestic_gross"].median())
gross_df['foreign_gross'] = gross_df['foreign_gross'].fillna(gross_df['foreign_gross'].median())
```

```
In [42]: # calculating 'total_gross' as the sum of 'domestic_gross' and 'foreign_gross'
gross_df['total_gross'] = gross_df['domestic_gross'] + gross_df['foreign_gross']
gross_df[['domestic_gross', 'foreign_gross', 'total_gross']].head()
```

Out[42]:

	domestic_gross	foreign_gross	total_gross
0	415000000.0	652000000.0	1.067000e+09
1	334200000.0	691300000.0	1.025500e+09
2	296000000.0	664300000.0	9.603000e+08
3	292600000.0	535700000.0	8.283000e+08
4	238700000.0	513900000.0	7.526000e+08

```
In [43]: #drop the rest with nulls
gross_df.dropna(inplace=True)
#confirming there are no null values
gross_df.isnull().sum().any()
```

Out[43]: False

```
In [44]: #Changing the numerical columns to currency for uniformity
#gross_df['domestic_gross'] = gross_df['domestic_gross'].apply(lambda x: f"${x:,.2f}")
#gross_df['foreign_gross'] = gross_df['foreign_gross'].apply(lambda x: f"${x:,.2f}")
#gross_df['total_gross'] = gross_df['total_gross'].apply(lambda x: f"${x:,.2f}")
```

```
In [45]: #adding the id as the first column
gross_df.insert(0, 'id', range(1, len(gross_df) + 1))
```

```
In [46]: gross_df.head()
```

Out[46]:

	id	title	domestic_gross	foreign_gross	year	total_gross
0	1	Toy Story 3	415000000.0	652000000.0	2010	1.067000e+09
1	2	Alice in Wonderland (2010)	334200000.0	691300000.0	2010	1.025500e+09
2	3	Harry Potter and the Deathly Hallows Part 1	296000000.0	664300000.0	2010	9.603000e+08
3	4	Inception	292600000.0	535700000.0	2010	8.283000e+08
4	5	Shrek Forever After	238700000.0	513900000.0	2010	7.526000e+08

### 3.3.2 Removing Outliers

```
In [47]: def remove_outliers(df, column_list):
    for column in column_list:
        # Calculate the first (Q1) and third (Q3) quartiles for the column
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)

        # Calculate the IQR
        IQR = Q3 - Q1

        # Define the lower and upper bounds for outliers
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        # Filter the dataframe to keep only the rows within the bounds
        df = df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]

    return df

# List of columns to remove outliers from
columns_to_check = ['foreign_gross', 'domestic_gross']

# Apply the function to remove outliers
gross_df = remove_outliers(gross_df, columns_to_check)

# Check the dataframe info to verify the results
gross_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2280 entries, 96 to 3386
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   id              2280 non-null   int64
 1   title           2280 non-null   object
 2   domestic_gross  2280 non-null   float64
 3   foreign_gross   2280 non-null   float64
 4   year            2280 non-null   int64
 5   total_gross     2280 non-null   float64
dtypes: float64(3), int64(2), object(1)
memory usage: 124.7+ KB
```

## 3.4 Cleaning for the tmdb\_df dataset

### 3.4.1 Dropping columns

```
In [48]: #Dropping the unnecessary columns
tmdb_df.drop(["vote_average", "vote_count", "genre_ids", "id", "original_language", "original_title", "overview", "release_date", "runtime", "spoken_languages", "status", "tagline", "title", "vote_average", "vote_count", "year"])
```

### 3.4.2 Checking for missing values

In [49]: *#Checking for the missing values*  
 tmdb\_df.isna().sum()

Out[49]: Unnamed: 0        0  
 popularity        0  
 release\_date      0  
 title             0  
 dtype: int64

In [50]: tmdb\_df.head()

Out[50]:

	Unnamed: 0	popularity	release_date	title
0	0	33.533	2010-11-19	Harry Potter and the Deathly Hallows: Part 1
1	1	28.734	2010-03-26	How to Train Your Dragon
2	2	28.515	2010-05-07	Iron Man 2
3	3	28.005	1995-11-22	Toy Story
4	4	27.920	2010-07-16	Inception

### 3.4.3 Changing Columns

In [51]: *# Renaming columns in tmdb\_df*  
 tmdb\_df.rename(columns={'Unnamed: 0': 'id'},inplace=True)  
 tmdb\_df.rename(columns={'release\_date': 'year'},inplace=True)

In [52]: tmdb\_df.head()

Out[52]:

	id	popularity	year	title
0	0	33.533	2010-11-19	Harry Potter and the Deathly Hallows: Part 1
1	1	28.734	2010-03-26	How to Train Your Dragon
2	2	28.515	2010-05-07	Iron Man 2
3	3	28.005	1995-11-22	Toy Story
4	4	27.920	2010-07-16	Inception



## 3.5 Cleaning for mbasic\_df dataset

### 3.5.1 Dropping columns

```
In [53]: #Dropping columns unnecessary for our analysis
mbasics_df.drop(["original_title", "runtime_minutes"], axis=1, inplace=True)
```

### 3.5.2 Checking for missing values

```
In [54]: #Checking missing values
mbasics_df.isnull().sum().any
```

```
Out[54]: <bound method NDFrame._add_numeric_operations.<locals>.any of movie_id
0
primary_title      0
start_year         0
genres            5408
dtype: int64>
```

```
In [55]: #Filling the missing values for the genres columns which is an object
for column in mbasics_df.select_dtypes(include=["object"]).columns:
    mbasics_df[column].fillna(mbasics_df[column].mode()[0], inplace=True)
```

```
In [56]: #Confirming for no missing values
mbasics_df.isnull().sum().any()
```

```
Out[56]: False
```

### 3.5.3 Renaming columns

```
In [57]: #Renaming the columns
mbasics_df.rename(columns={"movie_id": 'id'}, inplace=True)

mbasics_df.rename(columns={"start_year": 'year'}, inplace=True)

mbasics_df.rename(columns={"primary_title": 'title'}, inplace=True)
```

```
In [58]: mbasics_df.head()
```

```
Out[58]:
```

	id	title	year	genres
0	tt0063540	Sunghursh	2013	Action,Crime,Drama
1	tt0066787	One Day Before the Rainy Season	2019	Biography,Drama
2	tt0069049	The Other Side of the Wind	2018	Drama
3	tt0069204	Sabse Bada Sukh	2018	Comedy,Drama
4	tt0100275	The Wandering Soap Opera	2017	Comedy,Drama,Fantasy

## 3.6 Cleaning for the rating\_df dataset

### 3.6.1 Dropping unnecessary columns

```
In [59]: #Dropping unnecessary columns  
rating_df.drop(["numvotes"],axis=1,inplace=True)
```

### 3.6.2 Checking for missing values

```
In [60]: #Checking for missing values  
rating_df.isnull().sum().any()
```

Out[60]: False

```
In [61]: #Renaming columns  
rating_df.rename(columns={"movie_id":"id"},inplace=True)  
  
rating_df.rename(columns={"averagerating":"average_rating"},inplace=True)
```

```
In [62]: rating_df.head()
```

Out[62]:

	id	average_rating
0	tt10356526	8.3
1	tt10384606	8.9
2	tt1042974	6.4
3	tt1043726	4.2
4	tt1060240	6.5

### 3.6.3 Checking Duplicates

```
In [63]: #Checking for duplicates for movie_df  
movie_df.duplicated().sum()
```

Out[63]: 0

```
In [64]: #Checking for duplicates for gross_df  
gross_df.duplicated().sum()
```

Out[64]: 0

```
In [65]: #Checking for duplicates for tmdb_df  
tmdb_df.duplicated().sum()
```

Out[65]: 0

```
In [66]: #Checking for duplicates for mbasics_df
mbasics_df.duplicated().sum()
```

Out[66]: 0

```
In [67]: #Checking for duplicates for the rating_df dataset
rating_df.duplicated().sum()
```

Out[67]: 0

```
In [68]: directors_df.duplicated().sum()
```

Out[68]: 127639

```
In [69]: directors_df.drop_duplicates(inplace=True)
directors_df.duplicated().sum()
```

Out[69]: 0

### 3.6.4 Feature engineering

```
In [70]: # Split 'genre' into 'main_genre' and 'supporting_genre'
movie_df['main_genre'] = movie_df['genre'].str.split('|').str[0]
movie_df['supporting_genre'] = movie_df['genre'].str.split('|').apply(lambda x: x[1] if x[1] != '' else '')

# Preview the result
movie_df[['genre', 'main_genre', 'supporting_genre']].head()
```

Out[70]:

	genre	main_genre	supporting_genre
1	Drama Science Fiction and Fantasy	Drama	Science Fiction and Fantasy
6	Comedy	Comedy	
7	Drama	Drama	
8	Drama	Drama	
15	Comedy Drama Mystery and Suspense	Comedy	Drama Mystery and Suspense

```
In [71]: # Convert 'theater_date' and 'dvd_date' columns to datetime format
movie_df["year"] = pd.to_datetime(movie_df["year"]).dt.year

movie_df[['year']].head()
```

Out[71]:

	year
1	2012
6	2002
7	2006
8	2010
15	2010

```
In [72]: #From the new columns we can drop the columns further for easier analysis
movie_df.drop(["genre", "supporting_genre"], axis=1, inplace=True)

#Renaming the remaining column
movie_df.rename({"main_genre": "genre"}, axis=1, inplace=True)
```

```
In [73]: #Confirming changes
movie_df.head()
```

Out[73]:

	id	rating	director	year	box_office	genre
1	3	R	David Cronenberg	2012	600000	Drama
6	10	PG-13	Jake Kasdan	2002	41032915	Comedy
7	13	R	Ray Lawrence	2006	224114	Drama
8	14	R	Taylor Hackford	2010	134904	Drama
15	22	R	George Hickenlooper	2010	1039869	Comedy

```
In [74]: # convert 'release_date' to year
tmdb_df["year"] = pd.to_datetime(tmdb_df["year"]).dt.year
```

```
In [75]: # merging the movie_df and gross_df on 'id'
movie_basics_rating_df = pd.merge(mbasics_df, rating_df, on='id', how='inner')
```

In [76]: `movie_basics_rating_df.head()`

Out[76]:

	id	title	year	genres	average_rating
0	tt0063540	Sunghursh	2013	Action, Crime, Drama	7.0
1	tt0066787	One Day Before the Rainy Season	2019	Biography, Drama	7.2
2	tt0069049	The Other Side of the Wind	2018	Drama	6.9
3	tt0069204	Sabse Bada Sukh	2018	Comedy, Drama	6.1
4	tt0100275	The Wandering Soap Opera	2017	Comedy, Drama, Fantasy	6.5

In [77]: *#Merging the tm1 dataset("tmdb.movies.csv") to the original merged data set*  
*# First merge: Add tmdb1 to the existing merged DataFrame*  
`movie_basics_rating_df_final = pd.merge(movie_basics_rating_df, tmdb_df, on=`  
  
*# Second merge: Add bom1 to the updated merged DataFrame using the same column*  
`movie_basics_rating_df_final = pd.merge(movie_basics_rating_df, gross_df, on=`  
*# Check the result*  
`print(movie_basics_rating_df_final.head())`

	id_x	title	year_x	genres	average_rating
0	tt0315642	Wazir	2016	Action, Crime, Drama	7.1
1	tt0337692	On the Road	2012	Adventure, Drama, Romance	6.1
2	tt4339118	On the Road	2014	Drama	6.0
3	tt5647250	On the Road	2016	Drama	5.7
4	tt0376136	The Rum Diary	2011	Comedy, Drama	6.2

	id_y	domestic_gross	foreign_gross	year_y	total_gross
0	2569	1100000.0	18900000.0	2016	20000000.0
1	905	744000.0	8000000.0	2012	8744000.0
2	905	744000.0	8000000.0	2012	8744000.0
3	905	744000.0	8000000.0	2012	8744000.0
4	475	13100000.0	10800000.0	2011	23900000.0

In [78]: *#Changing the merged dataset to a dataframe*  
merged\_movies\_final\_df=pd.DataFrame(movie\_basics\_rating\_df\_final)  
merged\_movies\_final\_df

Out[78]:

	id_x	title	year_x	genres	average_rating	id_y
0	tt0315642	Wazir	2016	Action,Crime,Drama	7.1	2569
1	tt0337692	On the Road	2012	Adventure,Drama,Romance	6.1	905
2	tt4339118	On the Road	2014	Drama	6.0	905
3	tt5647250	On the Road	2016	Drama	5.7	905
4	tt0376136	The Rum Diary	2011	Comedy,Drama	6.2	475
...	...	...	...	...	...	...
1965	tt8290698	The Spy Gone North	2018	Drama	7.2	3303
1966	tt8331988	The Chambermaid	2018	Drama	7.1	2322

In [79]: *#Dropping unnecessary columns for the merged dataset*  
merged\_movies\_final\_df.drop(columns=["id\_x", "year\_x", "id\_y"],axis=1,inplace=True)


In [80]: *#Renaming columns*  
merged\_movies\_final\_df.rename({"year\_y":"year"},axis=1,inplace=True)

```
In [81]: #Confirming the changes
merged_movies_final_df
```

Out[81]:

	title	genres	average_rating	domestic_gross	foreign_gro
0	Wazir	Action, Crime, Drama	7.1	1100000.0	189000
1	On the Road	Adventure, Drama, Romance	6.1	744000.0	80000
2	On the Road	Drama	6.0	744000.0	80000
3	On the Road	Drama	5.7	744000.0	80000
4	The Rum Diary	Comedy, Drama	6.2	13100000.0	108000
...	...	...	...	...	...
1965	The Spy Gone North	Drama	7.2	501000.0	189000
1966	The Chambermaid	Drama	7.1	300.0	189000
1967	Helicopter Eela	Drama	5.4	72000.0	189000
1968	Last Letter	Drama, Romance	6.4	181000.0	189000
1969	Burn the Stage: The Movie	Documentary, Music	8.8	4200000.0	161000

1970 rows × 7 columns



### 3.6.5 Saving Dataset

```
In [82]: #Saving the merged dataset
merged_movies_final_df.to_csv("Datasets/merged_movies_clean.csv")
```


```
In [83]: #Saving the movie_df dataset
movie_df.to_csv("Datasets/movie_info_clean.csv")
```

## 4 Data Analysis

### 4.0.1 Analysis based on Genre Vs Rating

We will make a visualization for top 10 genres with the highest ratings

```
In [84]: genre_rating= merged_movies_final_df.groupby("genres")["average_rating"].sum
plt.figure(figsize=(5, 5))
sns.barplot(x=genre_rating.index, y=genre_rating.values,palette="Blues")
plt.xticks(rotation=45, ha='right')
plt.title('Genre vs Rating')
plt.show()
```



#### 4.0.1.1 Findings:


From the above analysis it is evident that the drama genre is the genre with highest rating with a rating higher than 2000.

#### 4.0.2 Analysis based on Genre Vs Sales

Here we will make visualization based on the top 10 highest selling genres based on amount each genre grossed.



```
In [85]: genre_sales= merged_movies_final_df.groupby("genres")["total_gross"].sum().\nplt.figure(figsize=(5, 5))\nsns.barplot(x=genre_sales.index, y=genre_sales.values,palette="Blues")\nplt.xticks(rotation=45, ha='right')\nplt.title('Genre vs Sales')\nplt.show()
```




#### 4.0.2.1 Findings:

Here we are able to view that the Adventure, Comedy and Sci-Fi genre has the highest sales. This may be brought by reasons such perfect direction, quality production and entertaining movies from this certain genre.

#### 4.0.3 Analysis based on Genre popularity

Here we will make visualizations based on genres with most popularity

```
In [86]: merged_movies_final_df["genres"].value_counts().head(5).plot(kind='bar',fig:
sns.set_palette('Blues')
plt.title("popular genres")
plt.ylabel("counts")
plt.xlabel('genre')
plt.xticks(rotation=45);
```



#### 4.0.3.1 Findings:

From the above analysis it is evident that drama genre is also the most popular genre. This may be due audience preferability, quality production and perfect direction from the directors associated with this particular genre

#### 4.0.4 Analysis on top selling movies with their popular genres

We can further our analysis based on best selling movies by visualizing with their associate genres

```
In [87]: #top movies based on genre popularity
top_three_movies = merged_movies_final_df["title"].value_counts().head(3).index


filtered_data = merged_movies_final_df[merged_movies_final_df["title"].isin(top_three_movies)]

top_genres=(
    filtered_data.groupby("title")["genres"]
    .value_counts()
    .groupby(level=0).nlargest(3)
    .reset_index(level=0,drop=True)
    .index.get_level_values(1)
)

filtered_data=filtered_data[filtered_data['genres'].isin(top_genres)]

plt.figure(figsize=(5, 5))#select figure size
sns.countplot(data=filtered_data, x='title', hue='genres')#selecting type of plot
plt.title('Top 3 Movies with Top 3 Genres Vs popularity')#title for the graph
plt.xlabel('Movie')#x-axis label
plt.ylabel('Popularity')#y-axis label
plt.legend(title='Genres')#legend title
plt.legend(fontsize='small') # You can use 'small', 'medium', 'large' or 'x-small'

# Alternatively, you can control the size of the Legend box:
plt.legend(handlelength=1, fontsize=8) #
plt.show()#visualize the graph
```



#### 4.0.4.1 Findings:

It is evident that the top selling movies are associated with the best rating and most popular genres such as the drama genre.


## 4.0.5 Analysis on best rated genre yearly

Since drama is the most popular genre we can create a visualization on how it has faired over the years

```
In [88]: #how genre with best rating has faired over the years(univariate analysis)
# Filter for the genre of interest
genre_of_interest = "Drama"
drama_df = merged_movies_final_df[merged_movies_final_df['genres'].str.contains(genre_of_interest)]

# Perform univariate analysis: Focus on average ratings over time
drama_yearly = drama_df.groupby('year')['average_rating'].mean()

# Plotting the univariate trend
plt.figure(figsize=(5, 5))
drama_yearly.plot(kind='line', marker='o', color='blue')
plt.title("Drama Genre Ratings Over the Years")
plt.xlabel("Year", fontsize=12)
plt.ylabel("Average Rating", fontsize=12)
plt.grid(visible=True, linestyle='--', alpha=0.6)
plt.show()
```




### 4.0.5.1 Findings:

The graph depicts the trend of average ratings for the drama genre from 2010 to 2018. Here's an analysis: The ratings exhibit fluctuations over the years rather than a consistent trend. The highest average rating occurs around 2014, reaching approximately 6.70. The lowest rating is observed in 2010, near 6.45. There is notable variability, with significant increases from 2010 to 2011 and a sharp rise to the peak in 2014. Post-2014, the ratings dip and rise again, peaking slightly in 2016 and 2017 before a small drop in 2018. From 2015 to 2016, the ratings appear more stable compared to previous years. This analysis suggests a variability in drama ratings, with an overall upward movement from the start to the peak, followed by a decline and stabilization towards the end of the period. This may be brought about by poor direction or production for this particular genre.

#### 4.0.6 Analysis based on top popular directors

We can further our analysis based on the top 5 directors

```
In [89]: #popular directors
director_popularity = movie_df['director'].value_counts().head(5)
plt.figure(figsize=(5,5))
sns.barplot(x=director_popularity.index,y=director_popularity.values)
plt.title("most popular directors")
plt.xlabel("director")
plt.ylabel("count")
plt.xticks(rotation=45);
```



##### 4.0.6.1 Findings:

From the above analysis it is evident that Clint Eastwood is the highest selling director. This may be brought by factors such as perfect direction, availability and good work ethics.

## 4.0.7 Analysis based on best selling directors

We can then make a visualization to view the top selling directors based on their box office.

Note: Number of tickets that are sold for a movie, as a measure of how popular and financially successful the movie or director is.

```
In [90]: #best selling directors


# Group by director and sum the box office sales
director_sales = movie_df.groupby("director")["box_office"].sum().sort_values

# Create the plot
plt.figure(figsize=(5, 5))
sns.barplot(x=director_sales.values, y=director_sales.index, palette="Blues")

# Add title and Labels
plt.xlabel('Box Office Sales')
plt.title('Directors vs Sales')

# Include x-tick labels (rotate for readability)
plt.xticks(rotation=60)

# Show the plot
plt.show()
```



### 4.0.7.1 Findings:

It is evident that Jay Roach made the highest sales based on box-office. This may be brought about by factors such as quality direction, good work ethics and professionalism.

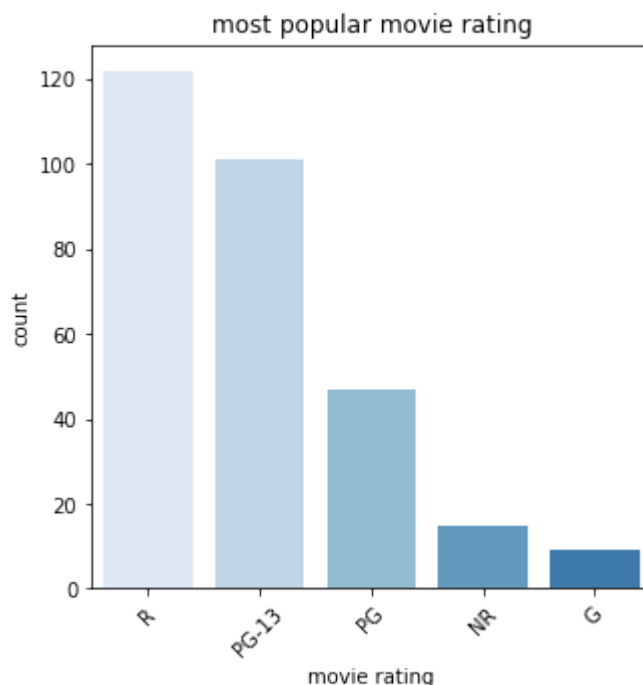
## 4.0.8 Analysis based on popular movie rating

We can further our analysis by movie rating popularity.

Note: In this visualization we are reviewing most popular assigned for each audience as follows:

1. R means restricted for audience under 18 years(adults)
2. PG-13 means restricted for audience under 13 years
3. PG means not restricted but requires parental guidance for audience less than 13 years
4. NR means not rated
5. G means for general audience meaning it is not restricted to any audience

```
In [91]: #popular movie rating
rating_popularity = movie_df['rating'].value_counts().head(5)
plt.figure(figsize=(5,5))
sns.barplot(x=rating_popularity.index,y=rating_popularity.values)
plt.title("most popular movie rating")
plt.xlabel("movie rating")
plt.ylabel("count")
plt.xticks(rotation=45);
```



#### 4.0.8.1 Findings

It is evident that the movie rating with the highest popularity is the R rated. This may be due to audience popularity who are adults.


#### 4.0.9 Analysis based on the best selling director yearly

Here we did analysis on the best selling director and he faired on yearly

```
In [92]: #univariate analysis based on the highest selling director
# Filter for the genre of interest
highest_director = "Jay Roach"
director_df = movie_df[movie_df['director'].str.contains(highest_director,

# Perform univariate analysis: Focus on average ratings over time
performance_yearly = director_df.groupby('year')['box_office'].mean()

# Plotting the univariate trend
plt.figure(figsize=(5, 5))
drama_yearly.plot(kind='line', marker='o', color='blue')
plt.title("Jay Roach's sales Over the Years")
plt.xlabel("Year", fontsize=12)
plt.ylabel("Sales", fontsize=12)
plt.grid(visible=True, linestyle='--', alpha=0.6)
plt.show()
```



#### 4.0.9.1 Findings:

The director sales have exhibit fluctuations over the years, with no consistent upward or downward trend. There are peaks and troughs at various points. The sales peaked significantly in 2014. There was another noticeable high point in 2017. The lowest sales occurred in 2010 and 2011. A decline in sales can also be observed in 2015 and 2018 after prior increases. Between 2012 and 2013 as well as 2015 to 2016, the sales remained relatively stable with minimal fluctuations. The variation in sales could be attributed to the performance of individual projects, changes in market dynamics, or external factors like competition or shifts in audience preferences.

## 4.1 Conclusions/ Results based on our analysis

From the analysis, it is clear that the Drama genre enjoys the highest ratings, but its box office success fluctuates over time. While the Adventure, Comedy, and Sci-Fi genres remain the highest-selling genres, this is likely due to their broad appeal and entertainment value. Directors like Clint Eastwood and Jay Roach contribute significantly to the sales, with their



strong reputations and professional standards. The ratings analysis also highlights the R rating's popularity, reflecting the tastes of adult audiences. Overall, the variability in movie sales and ratings over the years is likely driven by factors such as market dynamics, audience preferences, and the quality of direction and production.

## 4.2 Recommendations

1. We recommend the studio produce movies related to the drama genre
2. The studio to consider making movies associated with the Adventure, Comedy, Sci-Fi genres as the highest selling genres
3. Also recommend the studio to work with the director Clint Eastwood based on his popularity
4. The studio can also consider working with the highest selling director Jay Roach
5. Finally we recommend that the studio make adult films as it is the most popular among the audiences

## 5 Hypothesis Testing


- **95% Confidence Interval (CI):** The true value of the performance metric is expected to lie within this range with 95% confidence based on the data and model.
- **Alpha value ( $\alpha$ ):** The significance level is **0.05**, corresponding to a 95% confidence level, indicating that the likelihood of observing a value outside this range is 5%.

In [93]: merged\_movies\_final\_df.describe()

Out[93]:


	average_rating	domestic_gross	foreign_gross	year	total_gross
count	1970.000000	1.970000e+03	1.970000e+03	1970.000000	1.970000e+03
mean	6.460812	2.005813e+06	1.474976e+07	2014.114721	1.675557e+07
std	0.999176	3.573622e+06	8.693627e+06	2.366826	9.497521e+06
min	1.600000	1.000000e+02	6.000000e+02	2010.000000	1.080000e+04
25%	5.900000	5.950000e+04	7.000000e+06	2012.000000	1.102500e+07
50%	6.600000	3.280000e+05	1.890000e+07	2014.000000	1.895730e+07
75%	7.100000	2.000000e+06	1.890000e+07	2016.000000	1.970000e+07
max	9.200000	1.710000e+07	5.410000e+07	2018.000000	7.100000e+07

```
In [94]: merged_movies_final_df.average_rating.hist(color="blue");
```




```
In [95]: # apply log transformation to the 'average_rating' column to reduce skewness
merged_movies_final_df['log_average_rating'] = np.log1p(merged_movies_final_df['average_rating'])
merged_movies_final_df['log_average_rating'].hist(color="blue")
merged_movies_final_df['log_average_rating'].head()
```

```
Out[95]: 0    2.091864
1    1.960095
2    1.945910
3    1.902108
4    1.974081
Name: log_average_rating, dtype: float64
```



```
In [96]: # standardize the log-transformed 'average_rating' column
log_average_rating_mean = merged_movies_final_df['log_average_rating'].mean()
log_average_rating_std = merged_movies_final_df['log_average_rating'].std()
merged_movies_final_df['log_average_rating_zscore'] = (merged_movies_final_df['log_average_rating'] - log_average_rating_mean) / log_average_rating_std
# Plot a curve for the log-transformed and standardized average_ratings
plt.figure(figsize=(10, 6))
sns.kdeplot(merged_movies_final_df['log_average_rating_zscore'], shade=True)
plt.title('Density Curve of Log-Transformed Standardized average_ratings (Z-Scores)')
plt.xlabel('Z-Score', fontsize=14)
plt.ylabel('Density', fontsize=14)
plt.grid(True)
plt.show()
```



```
In [97]: merged_movies_final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1970 entries, 0 to 1969
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   title                 1970 non-null   object
 1   genres                1970 non-null   object
 2   average_rating        1970 non-null   float64
 3   domestic_gross        1970 non-null   float64
 4   foreign_gross         1970 non-null   float64
 5   year                 1970 non-null   int64
 6   total_gross           1970 non-null   float64
 7   log_average_rating    1970 non-null   float64
 8   log_average_rating_zscore 1970 non-null   float64
dtypes: float64(6), int64(1), object(2)
memory usage: 153.9+ KB
```

```
In [98]: print("Null hypothesis\n log_average_rating_mean = ", log_average_rating_mean)
print("Alternative hypothesis\n log_average_rating_mean > ", log_average_rating_mean)
```

```
Null hypothesis
log_average_rating_mean = 1.9994422593545513
Alternative hypothesis
log_average_rating_mean > 1.9994422593545513
```

```
In [99]: alpha = 0.05 # significance level

# extract the column data
log_average_rating_data = merged_movies_final_df['log_average_rating']

# sample statistics
n = np.random.randint(100, len(log_average_rating_data)) # sample size
sample_mean = np.mean(log_average_rating_data) # sample mean

# calculate the z-statistic
z_stat = (sample_mean - log_average_rating_mean) / (log_average_rating_std / np.sqrt(n))


# perform one-tailed z-test (alternative hypothesis: mean > mu)
p_value = 1 - norm.cdf(z_stat)

# output the results
print(f"Sample Mean: {sample_mean}")
print(f"Z-Statistic: {z_stat}")
print(f"P-Value: {p_value}")

# make a decision based on the p-value
if p_value < alpha:
    print("Reject the null hypothesis: The mean is significantly greater than the null hypothesis mean")
else:
    print("Fail to reject the null hypothesis: There is no significant evidence to suggest the mean is greater than the null hypothesis mean")
```


```
Sample Mean: 1.9994422593545513
Z-Statistic: 0.0
P-Value: 0.5
Fail to reject the null hypothesis: There is no significant evidence
that the mean is greater than 1.9994422593545513 at 95% confidence interval
```

```
In [100]: merged_movies_final_df.total_gross.hist(color="blue");
```




```
In [101]: # apply log transformation to the 'total_gross' column to reduce skewness
merged_movies_final_df['log_total_gross'] = np.log1p(merged_movies_final_df
merged_movies_final_df['log_total_gross'].hist(color="blue")
merged_movies_final_df['log_total_gross'].head()
```

```
Out[101]: 0    16.811243
1    15.983878
2    15.983878
3    15.983878
4    16.989389
Name: log_total_gross, dtype: float64
```



```
In [102]: # standardize the log-transformed 'total_gross' column
log_total_gross_mean = merged_movies_final_df['log_total_gross'].mean()
log_total_gross_std = merged_movies_final_df['log_total_gross'].std()
merged_movies_final_df['log_total_gross_zscore'] = (merged_movies_final_df[
# Plot a curve for the log-transformed and standardized total_grosss
plt.figure(figsize=(10, 6))
sns.kdeplot(merged_movies_final_df['log_total_gross_zscore'], shade=True, c
plt.title('Density Curve of Log-Transformed Standardized total_grosss (Z-Score)')
plt.xlabel('Z-Score', fontsize=14)
plt.ylabel('Density', fontsize=14)
plt.grid(True)
plt.show()
```



```
In [103]: print("Null hypothesis\n log_total_gross_mean = ", log_total_gross_mean)
print("Alternative hypothesis\n log_total_gross_mean > ", log_total_gross_mean)
```

Null hypothesis

log\_total\_gross\_mean = 16.269271504462154

Alternative hypothesis

log\_total\_gross\_mean > 16.269271504462154

```
In [104]: # extract the column data
log_total_gross_data = merged_movies_final_df['log_total_gross']

# sample statistics
n = np.random.randint(100, len(log_total_gross_data)) # sample size
sample_mean = np.mean(log_total_gross_data) # sample mean

# calculate the z-statistic
z_stat = (sample_mean - log_total_gross_mean) / (log_total_gross_std / np.s

# perform one-tailed z-test (alternative hypothesis: mean > mu)
p_value = 1 - norm.cdf(z_stat)

# output the results
print(f"Sample Mean: {sample_mean}")
print(f"Z-Statistic: {z_stat}")
print(f"P-Value: {p_value}")

# make a decision based on the p-value
if p_value < alpha:
    print("Reject the null hypothesis: The mean is significantly greater th
else:
    print("Fail to reject the null hypothesis: There is no significant evide
```

Sample Mean: 16.269271504462154

Z-Statistic: 0.0

P-Value: 0.5

Fail to reject the null hypothesis: There is no significant evidence  
that the mean is greater than 16.269271504462154 at 95% confidence interv  
al

## 6 Modeling

```
In [105]: merged_movies_final_df.head()
```

Out[105]:

	title	genres	average_rating	domestic_gross	foreign_gross	ye
0	Wazir	Action, Crime, Drama	7.1	1100000.0	18900000.0	
1	On the Road	Adventure, Drama, Romance	6.1	744000.0	8000000.0	
2	On the Road	Drama	6.0	744000.0	8000000.0	
3	On the Road	Drama	5.7	744000.0	8000000.0	
4	The Rum Diary	Comedy, Drama	6.2	13100000.0	10800000.0	

```
In [106]: #getting data for modelling
#merged_movies_modelling_df= merged_movies_final_df[["average_rating","domestic_gross","foreign_gross","total_gross"]]
merged_movies_modelling_df= merged_movies_final_df[["average_rating","domestic_gross","foreign_gross","total_gross"]]
```


```
In [107]: #feature Selection
selected_features = ["average_rating","domestic_gross","foreign_gross","total_gross"]
X = merged_movies_modelling_df[selected_features]
y = merged_movies_final_df['total_gross']
```

```
In [108]: # Correlation Analysis
correlation_matrix = merged_movies_final_df[selected_features].corr()
print("Correlation Matrix:\n", correlation_matrix)
```

Correlation Matrix:

	average_rating	domestic_gross	foreign_gross	total_gross
average_rating	1.000000	0.013149	0.044386	0.045577
domestic_gross	0.013149	1.000000	0.029821	0.403566
foreign_gross	0.044386	0.029821	1.000000	0.926578
total_gross	0.045577	0.403566	0.926578	1.000000

```
In [109]: plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```





```
In [110]: # Step 3: Modeling
# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ra
```

```
In [111]: model = sm.OLS(endog=y, exog=X)
result = model.fit()
result.summary()
```

Out[111]: OLS Regression Results

<b>Dep. Variable:</b>	total_gross	<b>R-squared (uncentered):</b>	1.000
<b>Model:</b>	OLS	<b>Adj. R-squared (uncentered):</b>	1.000
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	5.507e+32
<b>Date:</b>	Sat, 18 Jan 2025	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	20:42:23	<b>Log-Likelihood:</b>	32030.
<b>No. Observations:</b>	1970	<b>AIC:</b>	-6.405e+04
<b>Df Residuals:</b>	1967	<b>BIC:</b>	-6.404e+04
<b>Df Model:</b>	3		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>average_rating</b>	-3.893e-10	1.44e-10	-2.703	0.007	-6.72e-10	-1.07e-10
<b>domestic_gross</b>	0.3333	9.08e-17	3.67e+15	0.000	0.333	0.333
<b>foreign_gross</b>	0.3333	5.79e-17	5.76e+15	0.000	0.333	0.333
<b>total_gross</b>	0.6667	4.69e-17	1.42e+16	0.000	0.667	0.667

<b>Omnibus:</b>	181.928	<b>Durbin-Watson:</b>	0.453
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	603.703
<b>Skew:</b>	-0.440	<b>Prob(JB):</b>	8.08e-132
<b>Kurtosis:</b>	5.565	<b>Cond. No.</b>	1.79e+16


Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.

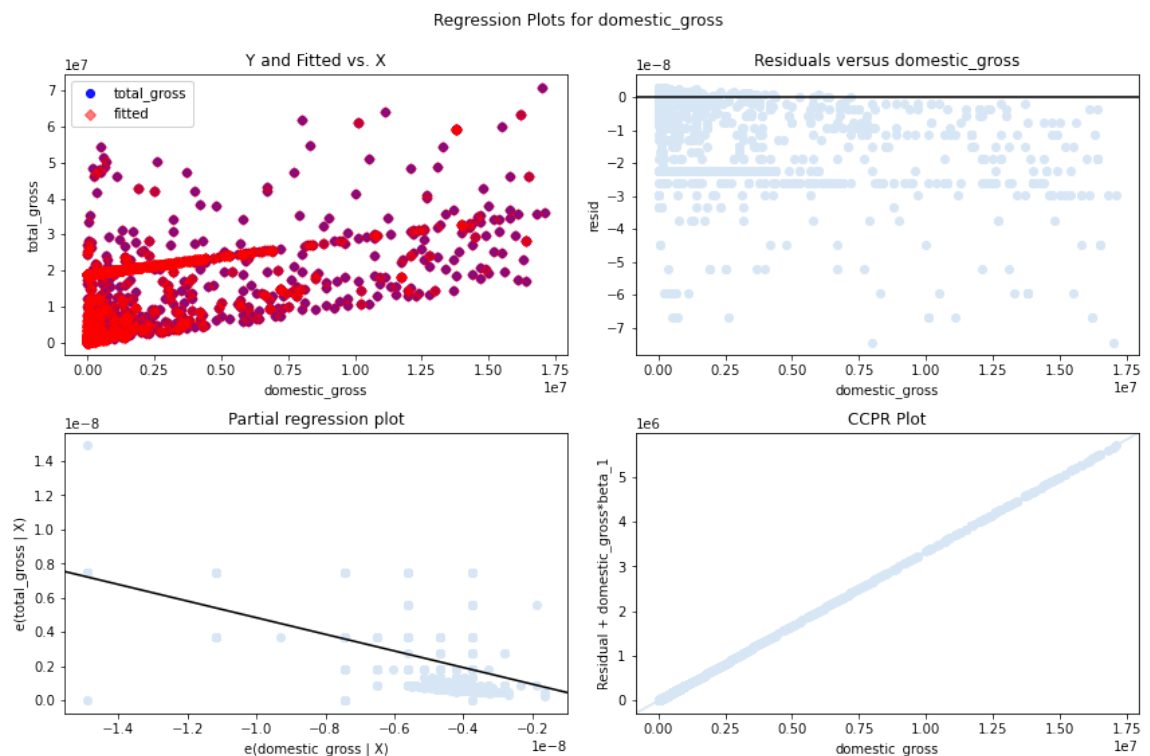
[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[3] The smallest eigenvalue is 4.09e-15. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.


In [112]: *#checking residual for average\_rating variable*  
`sm.graphics.plot_regress_exog(result, "average_rating", fig=plt.figure(figsize=(10, 10)))`



In [113]: *#checking residual for domestic\_gross variable*  
`sm.graphics.plot_regress_exog(result, "domestic_gross", fig=plt.figure(figsize=(10, 10)))`



```
In [114]: #checking residual for foreign_gross variable
sm.graphics.plot_regress_exog(result, "foreign_gross", fig=plt.figure(figsize=(10, 10)))
```



```
In [115]: merged_movies_final_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 1970 entries, 0 to 1969
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	title	1970 non-null	object
1	genres	1970 non-null	object
2	average_rating	1970 non-null	float64
3	domestic_gross	1970 non-null	float64
4	foreign_gross	1970 non-null	float64
5	year	1970 non-null	int64
6	total_gross	1970 non-null	float64
7	log_average_rating	1970 non-null	float64
8	log_average_rating_zscore	1970 non-null	float64
9	log_total_gross	1970 non-null	float64
10	log_total_gross_zscore	1970 non-null	float64

```
dtypes: float64(8), int64(1), object(2)
```

```
memory usage: 184.7+ KB
```

## 6.1 linear\_regression

### 6.1.1 Feature Preprocessing

In [116]:

```
# Label encode the 'genres' column (target variable)
label_encoder = LabelEncoder()
merged_movies_modelling_df['genres_encoded'] = label_encoder.fit_transform(
```

In [117]:

```
# Features (X) and Target (y)
X = merged_movies_modelling_df[['total_gross', 'average_rating']] # Include
y = merged_movies_modelling_df['genres_encoded']

# Split data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ran

# Step 2: Train the Logistic Regression Model (since we are predicting a cat
model = LogisticRegression(max_iter=100)
model.fit(X_train, y_train)
```

Out[117]: LogisticRegression()

## 7 Evaluation

In [118]:

```
# Calculate evaluation metrics
y_pred = model.predict(X_test) # Make predictions using the trained model

mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print(f"MAE: {mae}")
print(f"MSE: {mse}")
print(f"RMSE: {rmse}")
print(f"R²: {r2}")
```

```
MAE: 66.00507614213198
MSE: 8115.197969543147
RMSE: 90.08439359591176
R²: -0.38027741022485917
```

**Observation** The model's performance is suboptimal, as reflected by the following key observations: High Errors: The MAE (65.66), MSE (8101.98), and RMSE (90.01) indicate significant prediction errors on average. Negative  $R^2$  (-0.378): The negative  $R^2$  suggests the model performs worse than simply predicting the mean of the target variable. Model Improvements Needed: The model likely requires better feature engineering, data preprocessing, and possibly a different model approach to improve accuracy and fit.

```
In [119]: # Step 3: Modify the Random Prediction Function to Predict Genre

def predict_random_genre(model, merged_movies_modelling_df):
    # Randomly select a row from the dataset
    random_row = merged_movies_modelling_df.sample(1)

    # Extract the values for total_gross and average_rating
    total_gross = random_row['total_gross'].values[0]
    average_rating = random_row['average_rating'].values[0]

    # Make a prediction using the trained model
    genre_encoded_prediction = model.predict([[total_gross, average_rating]])

    # Convert the encoded prediction back to the original genre
    predicted_genre = label_encoder.inverse_transform([genre_encoded_prediction])

    # Return the random row and the predicted genre
    return random_row[['genres', 'total_gross', 'average_rating']].values, predicted_genre

# Example usage:
random_movie, predicted_genre = predict_random_genre(model, merged_movies_modelling_df)
print("Random Movie Data: ", random_movie)
print("Predicted Genre: ", predicted_genre)
```

```
Random Movie Data: [['Comedy,Drama' 19031000.0 6.4]]
Predicted Genre: Drama
```